

Execution report — step by step, with captured output

This document walks through an actual execution of the scanner and records what each step produced. Every block below is real terminal output, pasted unedited.

The reference run uses **demo mode**, because it is reproducible: the fixture pins its clock, so anyone running the same command on any machine on any day gets these exact numbers. **Step 8 is a real run against a live GitHub organization**, with the token permission probe that preceded it.

Environment: Windows 11, Python 3.12.10, httpx 0.28.1, Jinja2 3.1.6.

Step 1 — Environment setup

```
python -m venv .venv
```

```
.venv\Scripts\python.exe -m pip install -r requirements-dev.txt
```

Installs `httpx`, `Jinja2` and `pytest`. `Chart.js` is already vendored in `vendor/chart.umd.min.js` so the dashboard needs no network at view time.

Step 2 — Test suite

```
python -m pytest -q
```

```
..... [ 52%]
..... [100%]
156 passed in 0.32s
```

156 tests, 0.32 seconds. They cover:

FILE	WHAT IT PINS
tests/test_score.py (80)	The scoring arithmetic against hand-computed values, decay behaviour, risk ordering, every NEVER-FLAG rule, and the remediation advice for team-inherited and archived grants
tests/test_client.py (33)	Rate-limit pauses, <code>Retry-After</code> , 409/403/404 handling, ETag 304s, cursor pagination, GraphQL error types
tests/test_pipeline.py (32)	The whole fixture org end-to-end: who is flagged, in exactly what order, who is excluded and why, plus database idempotency
tests/test_access.py (11)	A refused collaborator listing is reported, never rendered as "nobody has access"; org base permission is not team-inherited

The same suite runs in CI on Python 3.11 and 3.12, followed by a full `--demo` run whose findings are asserted against the numbers published in Step 3 below — so this report cannot silently drift from the code.

The scoring tests assert exact arithmetic, not plausibility — for example that `activity_score(commits=10)` equals `7.1936858`, which is `3.0 × ln(11)` computed independently.

Step 3 — The scan

```
python main.py --demo

13:31:23 INFO scanner: Demo mode: acme-demo from fixturesdemo_org.json, clock pinned to 2026-08-
13:31:23 INFO report: Dashboard written to outdashboard.html (256 KB)

Run #1 - acme-demo (demo)
-----
Repositories scanned      5 (1 skipped)
Advisories found         10 (8 still open)
Access suggestions       10
Excluded from review     10

Highest risk:
  3.00 alex-departed - admin on acme-demo/payments-api
  3.00 sam-stale - admin on acme-demo/legacy-etl
  2.00 maya-platform - maintain on acme-demo/legacy-etl (team-inherited)
  1.50 alex-departed - write on acme-demo/web-frontend (team-inherited)
  1.50 maya-platform - write on acme-demo/payments-api (team-inherited)

Dashboard: outdashboard.html
```

What happened, and why each number is what it is

6 repositories in the org, 5 scanned. `vendor-sdk-fork` is a fork and was skipped (`SKIP_FORKS` defaults on) — recorded as a row with a reason, not silently dropped. `legacy-etl` is archived and **was** scanned: stale admin access on an archived repo is still live access.

10 advisories, 8 open. Includes 2 critical and 2 high still open. The oldest unresolved is a Django SQL-injection advisory open for **347 days**.

10 access suggestions from 25 collaborator grants. The two highest are dormant admins scoring 0.00, so their risk equals the raw admin weight of 3.00.

Three of the top five are team-inherited and are labelled as such, with an action column explaining that removing them means changing team membership.

Who was *not* flagged, and why that is the interesting part

PERSON	SITUATION	OUTCOME
priya-staff	2 commits, 21 reviews on payments-api	Not flagged , score 13.25 — a commit-count model would have flagged her

PERSON	SITUATION	OUTCOME
lena-active	62 commits, 9 reviews	Not flagged, score 26.29
ravi-owner	Admin, 1 commit in 6 months	Excluded — org owner, never flagged regardless of activity
dependabot[bot]	Write access, 24 commits	Excluded — bot
release-bot	Write access, 15 commits	Excluded — bot (typed Bot , no [bot] suffix)
audit-svc	Read access, zero activity	Excluded — on the configured allowlist
everyone on ml-sandbox	Repo created 28 days ago	Excluded — no time to build a history
tom-triage on infra-scripts	Write, zero activity	Excluded — repo has only one contributor

priya-staff is the case the scoring model exists to protect, and a dedicated test fails if a future change flags her.

Step 4 — Idempotency check

Running the identical command a second time against the same database:

```
python main.py --demo
```

```
### row counts after two runs
runs          2
repos         6
advisories    10
collaborators 25
contributions 25
scores        25
exclusions    14
run_progress   40
```

Two runs, two runs rows — and every fact table unchanged. 6 repos, not 12. 10 advisories, not 20. 25 scores, not 50. The facts were updated in place and re-stamped with the newer run_id ; only the run ledger and the per-run progress log grew, which is exactly what should grow.

This is the idempotency requirement demonstrated rather than asserted, and test_pipeline.py::test_seeding_twice_does_not_duplicate_rows enforces it on every test run.

Step 5 — Re-rendering without touching the API

```
python main.py --report-only
```

```
2.00 maya-platform - maintain on acme-demo/legacy-etl (team-inherited)
1.50 alex-departed - write on acme-demo/web-frontend (team-inherited)
1.50 maya-platform - write on acme-demo/payments-api (team-inherited)
```

Dashboard: out\dashboard.html

Rebuilds the dashboard from SQLite with zero network calls — useful for iterating on the report, and for regenerating a view of a scan captured days earlier.

Step 6 — Machine-readable output

```
python main.py --demo --json
```

```
{
  "run_id": 2,
  "org": "acme-demo",
  "mode": "demo",
  "started_at": "2026-08-25T07:58:24+00:00",
  "finished_at": "2026-08-25T07:58:24+00:00",
  "status": "completed",
  "repos_scanned": 5,
  "repos_skipped": 1,
  "repos_errorred": 0,
  "repos_excluded": 4,
  "advisories_found": 10,
  "advisories_open": 8,
  "suggestions_made": 10,
  "members_excluded": 10,
  "collaborators_seen": 25,
  "top_suggestions": [
    {"login": "alex-departed", "repo": "acme-demo/payments-api",
     "permission": "admin", "risk": 3.0, "score": 0.0, "team_inherited": false},
    {"login": "sam-stale", "repo": "acme-demo/legacy-etl",
     "permission": "admin", "risk": 3.0, "score": 0.0, "team_inherited": false},
    {"login": "maya-platform", "repo": "acme-demo/legacy-etl",
     "permission": "maintain", "risk": 2.0, "score": 0.0, "team_inherited": true}
  ],
  "dashboard": "out/dashboard.html",
  "database": "data/scanner.sqlite3",
  "api": {}
}
```

Logs go to stderr, so stdout stays clean and pipeable:

```
python main.py --demo --json > summary.json
```


"api" is empty in demo mode because no requests were made. On a live run it carries request count, cache hits, retries and rate-limit pauses.

Step 7 — The dashboard

`out/dashboard.html` — one self-contained file, 256 KB with Chart.js inlined. Open it directly in a browser; no server required. A committed copy lives at [docs/demo_dashboard.html](#), and a printed version at [docs/pdf/dashboard.pdf](#).


The screenshots below are generated by `python tools/make_docs.py`, so they cannot drift from the dashboard they document.

Header — the answer before the data



A security lead reads three sentences and knows the state of the org. Everything below is the evidence for those three sentences.

Section 1 — Security advisories



Section 2 — Who has access, and what they did with it

acme-demo — security & access hygiene

Generated 25 Aug 2026, 12:00 UTC · Activity window: last 180 days · 5 repositories scanned, 1 skipped · Run #2

2 Who has access, and what they did with it

One row per person per repository. **Score** combines commits, reviews and pull requests over the last 180 days, with older work counting for less. A score below **5.0** triggers a review suggestion. Repositories are ordered by how many suggestions they produced. Click any repository to expand it.

payments-apiprivate3 open alerts5 to review10 with access

Created 04 Mar 2021 · last push 24 Aug 2026 · 99 commits in window from 5 identified contributors · 7 commits could not be attributed (commit email not linked to any GitHub account — deliberately not guessed) · [open on GitHub](#)

MEMBER	HOW THEY HAVE ACCESS	COMMITTS	PRS REVIEWED	PRS MERGED	LAST ACTIVITY	SCORE	STATUS
lena-active	Direct (Write)	62	9	14	24 Aug 2026 (0 days ago)	26.29	active
dependabot[bot]	Direct (Write)	24	0	22	23 Aug 2026 (2 days ago)	18.97	excluded — Bot or app account
priya-staff	Direct (Maintain)	2	21	1	20 Aug 2026 (5 days ago)	13.25	active
ravi-owner	Direct (Admin)	1	2	0	11 Jun 2026 (75 days ago)	3.61	excluded — Organization owner
dev-vendor	Outside collaborator (Write)	3	0	1	05 Mar 2026 (172 days ago)	3.21	review access
alex-departed	Direct (Admin)	0	0	0	17 Apr 2025 (495 days ago) before the window	0.00	review access
audit-svc	Direct (Read)	0	0	0	never	0.00	excluded — On the configured allowlist
maya-platform	Team: Platform Engineering (Write) team-inherited	0	0	0	30 Aug 2025 (360 days ago) before the window	0.00	review access
sam-stale	Direct (Write)	0	0	0	14 Nov 2025 (283 days ago) before the window	0.00	review access
tom-triage	Direct (Triage)	0	0	0	never	0.00	review access

One row per person per repo. `lena-active` scores 26.29; `priya-staff` scores 13.25 on 2 commits and 21 reviews — the case a commit-count model gets wrong. Excluded people stay visible with their reason attached rather than vanishing. The panel header also reports the **7 unattributed commits** on this repository.

Section 3 — Suggested access removals

acme-demo — security & access hygiene

Generated 25 Aug 2026, 12:00 UTC · Activity window: last 180 days · 5 repositories scanned, 1 skipped · Run #2

3

Suggested access removals

Ranked by **risk** = how much power the grant carries, divided by how much the person is using it. A dormant admin ranks above a dormant reader. **Nothing here has been changed** — this tool never revokes access. Every row is a recommendation for a human to confirm.

#	MEMBER	REPOSITORY	PERMISSION	LAST COMMIT	LAST REVIEW	SCORE	RISK	WHY FLAGGED	WHAT TO DO
1	alex-departed	payments-api	Admin	17 Apr 2025 (495 days ago) before the window	never	0.00	3.00	Holds admin access but has no recorded commits, reviews or pull requests in the last 180 days. Score 0.00 is below the 5.0 threshold.	Direct repository grant can be revoked on this repo alone.
2	sam-stale	legacy-etl archived	Admin	02 Jun 2024 (814 days ago) before the window	never	0.00	3.00	Holds admin access but has no recorded commits, reviews or pull requests in the last 180 days. Score 0.00 is below the 5.0 threshold.	Direct repository grant can be revoked on this repo alone. The repository is archived, and GitHub refuses collaborator changes while it is - remove the grant at the org or team level, or unarchive, revoke, and re-archive.
3	maya-platform	legacy-etl archived	Maintain via Data Engineering	08 Nov 2023 (1020 days ago) before the window	never	0.00	2.00	Holds maintain access but has no recorded commits, reviews or pull requests in the last 180 days. Score 0.00 is below the 5.0 threshold.	Access is inherited from Data Engineering. Removing it means changing team membership, which affects every repo that team can reach - review at the team level, not here. The repository is archived, and GitHub refuses collaborator changes while it is - remove the grant at the org or team level, or unarchive, revoke, and re-archive.
4	alex-departed	web-frontend	Write via Platform Engineering	21 Jun 2025 (429 days ago) before the window	never	0.00	1.50	Holds write access but has no recorded commits, reviews or pull requests in the last 180 days. Score 0.00 is below the 5.0 threshold.	Access is inherited from Platform Engineering. Removing it means changing team membership, which affects every repo that team can reach - review at the team level, not here.
5	maya-platform	payments-api	Write via Platform Engineering	30 Aug 2025 (360 days ago) before the window	never	0.00	1.50	Holds write access but has no recorded commits, reviews or pull requests in the last 180	Access is inherited from Platform Engineering. Removing it means changing team

Ranked by risk. Every row carries last commit, last review, permission, how the access was granted, why it was flagged, and what the reviewer would have to do about it. Team-inherited grants say "review at the team level"; grants on an archived repo add that GitHub refuses collaborator changes while a repo is archived, so the fix is at org or team level.

Section 4 — Who was excluded, and why

acme-demo — security & access hygiene

Generated 25 Aug 2026, 12:00 UTC · Activity window: last 180 days · 5 repositories scanned, 1 skipped · Run #2

4 Who was excluded, and why

These people and repositories were deliberately kept out of the suggestions above. This panel exists so the list you just read can be trusted: if you cannot see who was skipped, you have no way to know whether the tool is hiding something.

People never suggested for removal (10)

On the configured allowlist (1)

audit-svcacme-demo/payments-apiRead

Bot or app account (2)

dependabot[bot]acme-demo/payments-apiWrite

release-botacme-demo/web-frontendWrite

Organization owner (2)

ravi-owneracme-demo/legacy-etlAdmin

ravi-owneracme-demo/payments-apiAdmin

Repository created within the last 180 days (3)

alex-departedacme-demo/ml-sandboxWrite

lena-activeacme-demo/ml-sandboxWrite

noor-newacme-demo/ml-sandboxAdmin

Repository has only one contributor (2)

kai-soloacme-demo/infra-scriptsAdmin

tom-triageacme-demo/infra-scriptsWrite

Repositories and how they were treated

Repository is archived (1)

legacy-etlArchived, but still scanned - access on an archived repo is still live access, and anyone holding it can unarchive the repo.

Repository is a fork (1)

vendor-sdlk-forkfork (SKIP_FORKS is on)

Repository created within the last 180 days (1)

ml-sandboxCreated 28 days ago, inside the 180-day window. Nobody has had time to build a history here yet.

Repository has only one contributor (1)

infra-scriptsOnly 1 contributor in the window. With no one to compare against, a low score says nothing useful.

Unattributed commits (22)

Commits whose author email is not linked to any GitHub account. They are counted and reported but never assigned to a person, because guessing would credit — or fail to credit — the wrong someone.

Org owners, bots, allowlisted accounts, and repositories too new or too small to judge — each named, grouped by reason. This panel is what makes the list in section 3 trustworthy: a reader who cannot see who was skipped has no reason to believe the people who were not.

Step 8 — Live run against a real organization

The fixture org above is the illustrative case. This is the same code against a real GitHub organization (VoidAlgo), and it is here because running against real data found four bugs that no fixture would have caught.

Token permission probe

Every endpoint the tool depends on, probed individually before scanning:

ENDPOINT	RESULT
/rate_limit	OK — REST 5000/5000, GraphQL 5000/5000
/orgs/{org}/members?role=admin	OK — 2 owners
/orgs/{org}/teams	OK — 0 teams
/orgs/{org}/repos	OK — 4 repositories
/orgs/{org}/dependabot/alerts	OK — 32 alerts in a single call

ENDPOINT	RESULT
/repos/{r}/collaborators (×3 affiliations)	OK on all 4
/repos/{r}/commits , /pulls , GraphQL	OK — 231 commits, 34 PRs

The org-wide advisory endpoint is the efficiency decision paying off in practice: **one paginated call returned all 32 alerts** rather than four per-repo calls, and that ratio is what keeps a 100-repo org affordable.

The scan

```
INFO scanner: Started run #1 for VoidAlgo
INFO scanner: Token OK. REST quota 5000/5000, GraphQL 5000/5000
INFO scan: Org VoidAlgo has 2 owner(s)
INFO scan: Org base permission is 'read': every member holds read on every repo
INFO scan: Found 7 repos in VoidAlgo; scanning 7, skipping 0
INFO scan: Org-level Dependabot endpoint returned 32 alerts across 1 repos
INFO scanner: [1/7] VoidAlgo/Agentic-CMO
INFO scanner: [2/7] VoidAlgo/Auto_Job_Applying_Agent
INFO scanner: [3/7] VoidAlgo/demo-repository
INFO scanner: [4/7] VoidAlgo/github_pull_toslack
INFO scanner: [5/7] VoidAlgo/pseudoquant
INFO scanner: [6/7] VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant
INFO scanner: [7/7] VoidAlgo/voice_mvp_dupe

Run #1 - VoidAlgo (live)
-----
Repositories scanned      7
Advisories found          32 (32 still open)
Access suggestions        2
Excluded from review      19
API                       51 requests, 0 served from cache, 0 rate-limit pauses

Highest risk:
  0.97 fatbatman85 - admin on VoidAlgo/voice_mvp_dupe
  0.49 fatbatman85 - write on VoidAlgo/Agentic-CMO
```

Advisory findings

32 open Dependabot alerts, 15 of them critical or high, all concentrated in a single repository:

SEVERITY	OPEN
Critical	1
High	14
Medium	11
Low	6

A sample, straight from the run:

SEVERITY	PACKAGE	ADVISORY	SUMMARY
Critical	torch	GHSA-53q9-r3pm-6pq6	PyTorch remote code execution
High	black	GHSA-3936-cmfr-pm3m	Arbitrary file writes from unsanitized input
High	pillow	GHSA-45hq-cxwh-f6vc	<code>Image.new()</code> bypasses the decompression-bomb check
Medium	Pillow	GHSA-4x4j-2g7c-83w6	<code>WindowsViewer.get_command()</code> OS command injection

One repository accounts for every open alert in the organization, which is exactly the "most affected repositories" signal the dashboard exists to surface.

Access findings — two suggestions, and the reasoning behind each

#	MEMBER	REPOSITORY	PERMISSION	SCORE	RISK
1	fatbatman85	voice_mvp_dupe	Admin	2.08	0.97
2	fatbatman85	Agentic-CMO	Write	2.08	0.49

Holds admin access but has contributed only 1 commit, 0 PRs reviewed, 0 PRs merged in the last 180 days. Score 2.08 is below the 5.0 threshold.

Both are genuine low-contribution findings: a single commit against a threshold calibrated at roughly four. The admin grant ranks first because risk divides permission weight by contribution — same score, twice the consequence.


What was *not* flagged, and why that matters more

SITUATION	OUTCOME
fatbatman85 holds admin on Auto_Job_Applying_Agent and made 3 commits	not flagged — active people keep their access
Makilesh, clashonkishy — admin everywhere	excluded — organization owners, a hard rule
fatbatman85 read on 4 repos	excluded — org base permission, no repo grant to revoke
Copilot authored commits on Sensitive_Data_Detection...	excluded — holds no access, nothing to remove
demo-repository, github_pull_toslack	excluded — one contributor, nothing to compare against
pseudoquant, Sensitive_Data_Detection...	excluded — created inside the 180-day window

19 access grants were assessed and held back, 2 were suggested. A tool that flagged the active admin, or the owners, or the person whose read access comes from an org-wide setting, would be worse than useless — the reviewer would stop trusting the list on the first false positive.

The live dashboard

The same four sections, rendered from the real organization. These are produced by `python tools/make_docs.py`, from the committed `docs/live/dashboard.html`, so they cannot drift from what the scan actually produced. A printed copy is at `docs/pdf/live_dashboard.pdf`.




Three sentences, then the numbers: 15 critical or high vulnerabilities open, 2 stale access grants, 21 people-repository pairs assessed.

1 Security advisories


Every Dependabot alert across the organization, grouped by how serious it is and whether anyone has dealt with it. **Open** means still exploitable today.

Open alerts by severity



SEVERITY	OPEN	FIXED	DISMISSED
Critical	1	0	0
High	14	0	0
Medium	11	0	0
Low	6	0	0

Status of all alerts



32 alerts total.

Most affected repositories



Oldest unresolved alerts

Sorted by how long they have been open. These are the ones that have been ignored the longest.

AGE	SEVERITY	PACKAGE	REPOSITORY	ADVISORY	SUMMARY
0d	Low	torch pip	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	GHSA-3749-ghw9-m3mg CVE-2025-2953	PyTorch susceptible to local Denial of Service
0d	High	black pip	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	GHSA-3936-cmfr-pm3m CVE-2026-32274	Black: Arbitrary file writes from unsanitized user input in cache file name
0d	High	pillow pip	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	GHSA-45hq-cxwh-f6vc CVE-2026-55379	Pillow 'BdfFontFile': 'image.new()' called without '_decompression_bomb_check()' — bomb protection bypass via font loading
0d	Medium	Pillow pip	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	GHSA-4x4j-2q7c-83w6 CVE-2026-55798	Pillow: WindowsViewer.get_command() OS command injection via unescaped shell path
0d	Critical	torch pip	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	GHSA-53q9-r3om-6pg6 CVE-2025-32434	PyTorch: 'torch.load' with 'weights_only=True' leads to remote code execution
0d	Medium	Markdown pip	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	GHSA-5wmx-573v-2qwg CVE-2025-69534	Python-Markdown has an Uncaught Exception
0d	High	pillow pip	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	GHSA-5x9d-69rx-g8h2 CVE-2026-54060	Pillow: 'FontFile.compile()': 'image.new()' called without '_decompression_bomb_check()'
0d	High	pillow pip	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	GHSA-62e4-gmf7-7g93 CVE-2026-54058	Pillow: Out-of-bounds read via attacker-controlled row stride on Pillow's mmap path (Mcidias AREA files)

32 open alerts, all in one repository — real CVEs against `torch`, `pillow` and `black`. The "most affected repositories" chart collapses to a single bar, which is itself the finding.

VoidAlgo — security & access hygiene

Generated 27 Aug 2026, 19:55 UTC · Activity window: last 180 days · 7 repositories scanned · Run #1

2 Who has access, and what they did with it

One row per person per repository. **Score** combines commits, reviews and pull requests over the last 180 days, with older work counting for less. A score below **5.0** triggers a review suggestion. Repositories are ordered by how many suggestions they produced. Click any repository to expand it.

Agentic-CMO

private

1 to review

3 with access

Created 27 Jan 2026 · last push 27 Aug 2026 · 2 commits in window from 2 identified contributors · [open on GitHub](#)

MEMBER	HOW THEY HAVE ACCESS	COMMITTS	PRS REVIEWED	PRS MERGED	LAST ACTIVITY	SCORE	STATUS
Makilesh	Direct (Admin)	1	0	0	27 Aug 2026 (0 days ago)	2.08	excluded — Organization owner
fatbatman85	Direct (Write)	1	0	0	27 Aug 2026 (0 days ago)	2.08	review access
clashonkishy	Organization owner (Admin)	0	0	0	never	0.00	excluded — Organization owner

Seven repositories, each expandable. Note `Auto_Job_Applying_Agent` : the same person who is flagged elsewhere holds **admin** here and is **not** flagged, because they actually contributed.

VoidAlgo — security & access hygiene

Generated 27 Aug 2026, 19:55 UTC · Activity window: last 180 days · 7 repositories scanned · Run #1

3 Suggested access removals

Ranked by **risk** = how much power the grant carries, divided by how much the person is using it. A dormant admin ranks above a dormant reader. **Nothing here has been changed** — this tool never revokes access. Every row is a recommendation for a human to confirm.

#	MEMBER	REPOSITORY	PERMISSION	LAST COMMIT	LAST REVIEW	SCORE	RISK	WHY FLAGGED	WHAT TO DO
1	fatbatman85	voice_mvp_dupe	Admin	27 Aug 2026 (0 days ago)	never	2.08	0.97	Holds admin access but has contributed only 1 commit, 0 PRs reviewed, 0 PRs merged in the last 180 days. Score 2.08 is below the 5.0 threshold.	Direct repository grant; can be revoked on this repo alone.
2	fatbatman85	Agentic-CMO	Write	27 Aug 2026 (0 days ago)	never	2.08	0.49	Holds write access but has contributed only 1 commit, 0 PRs reviewed, 0 PRs merged in the last 180 days. Score 2.08 is below the 5.0 threshold.	Direct repository grant; can be revoked on this repo alone.

Two suggestions, ranked by risk, each carrying its evidence and its remediation.

VoidAlgo — security & access hygiene

Generated 27 Aug 2026, 19:55 UTC · Activity window: last 180 days · 7 repositories scanned · Run #1

4 Who was excluded, and why

These people and repositories were deliberately kept out of the suggestions above. This panel exists so the list you just read can be trusted: if you cannot see who was skipped, you have no way to know whether the tool is hiding something.

People never suggested for removal (19)

Contributed here, but holds no current access - nothing to remove (1)

Copilot VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant

Organization owner (14)

Makilesh	VoidAlgo/Agentic-CMO	Admin
Makilesh	VoidAlgo/Auto_Job_Applying_Agent	Admin
Makilesh	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	Admin
Makilesh	VoidAlgo/demo-repository	Admin
Makilesh	VoidAlgo/github_pull_toslack	Admin
Makilesh	VoidAlgo/pseudoquant	Admin
Makilesh	VoidAlgo/voice_mvp_dupe	Admin
clashonkishy	VoidAlgo/Agentic-CMO	Admin
clashonkishy	VoidAlgo/Auto_Job_Applying_Agent	Admin
clashonkishy	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	Admin
clashonkishy	VoidAlgo/demo-repository	Admin
clashonkishy	VoidAlgo/github_pull_toslack	Admin
clashonkishy	VoidAlgo/pseudoquant	Admin
clashonkishy	VoidAlgo/voice_mvp_dupe	Admin

Repository created within the last 180 days (4)

fatbatman85	VoidAlgo/Sensitive_Data_Detection-Compliance_Assistant	Read
fatbatman85	VoidAlgo/demo-repository	Read
fatbatman85	VoidAlgo/github_pull_toslack	Read
fatbatman85	VoidAlgo/pseudoquant	Read

Repositories and how they were treated

Repository created within the last 180 days (4)

Sensitive_Data_Detection-Compliance_Assistant	Created 0 days ago, inside the 180-day window. Nobody has had time to build a history here yet.
demo-repository	Created 33 days ago, inside the 180-day window. Nobody has had time to build a history here yet.
github_pull_toslack	Created 0 days ago, inside the 180-day window. Nobody has had time to build a history here yet.
pseudoquant	Created 33 days ago, inside the 180-day window. Nobody has had time to build a history here yet.

Unattributed commits (6)

Commits whose author email is not linked to any GitHub account. They are counted and reported but never assigned to a person, because guessing would credit — or fail to credit — the wrong someone.

Nineteen grants held back, grouped by reason — organization owners, org base permission, contributors with no access, repositories too new or too small. This panel is what makes the two suggestions above believable.

A limitation this run exposed

`voice_mvp_dupe` was, before the commits above, a repository with **three admins and no commits in six months**. That is arguably the strongest possible stale-access signal — an abandoned repository — and the single-contributor rule suppressed it, because the rule was written for "only one person to compare against" and treats zero contributors the same as one. Those are different situations: one is thin evidence, the other is the evidence. Splitting them is the first change I would make to the exclusion rules.

Four bugs the live run found that fixtures never would

1. **A refused collaborator listing rendered as "nobody has access."** A token without `Administration:read` produced empty access lists that the report stated with full confidence. Now `AccessSnapshot.complete` separates *empty* from *unreadable*, and the dashboard leads with a red banner.
2. **Organization base permission mislabelled as team-inherited.** `default_repository_permission = "read"` puts every member into the `all` listing with no team involved. The remediation is one org-wide setting, not a team change, so it is now its own access path.
3. **Contributors with no access suggested for removal.** `Copilot` authored commits on a repository where it holds no permission, scored 2.08, and would have been recommended for the removal of access it never had.
4. **Advisory counts silently zeroed.** Two stages write `repo_stats`; the contribution pass overwrote the advisory pass's `open_advisories` with its default, so a repository with 32 stored alerts displayed "0 open alerts". The upsert is now a partial update.

All four became code changes with tests. Section 8 of [DESIGN_NOTES.md](#) records the reasoning.

Idempotency and caching, verified live

A second identical run:

```
"api": { "requests": 13, "cache_hits": 8, "retries": 0, "rate_limit_sleeps": 0 }
```

8 of 13 requests served from stored ETags — 304 responses, which do not count against the API quota. Row counts after two full runs: `runs 2`, every fact table unchanged. Idempotency and conditional caching confirmed against the real API, not against mocks.

What is mocked, stated plainly

Only the GitHub API responses, and only in `--demo`. Private Dependabot advisories require an organization with paid security features and genuinely vulnerable dependencies, which is not reproducible for a review — the brief explicitly permits mocking exactly this.

The fixtures are shaped like real API payloads and are fed through the **same** `db.upsert_advisory`, `contrib.AccessEntry` and `score.assess_repo` functions the live scan uses. There is no demo-only parsing or scoring path. What is mocked is the network; everything above it is the production code.