

GitHub Org Security & Access-Hygiene Automation — submission summary

Makilesh · SDE/Automation intern task · github.com/Makilesh/Github_Org_Security

Scans every repository in a GitHub organization, reports its security advisories, measures how much each collaborator actually contributes to each repo, and **suggests** access removals. It never revokes anything.

Contribution scoring, and why

```
activity = 3.0·ln(1+commits) + 2.5·ln(1+reviews) + 2.0·ln(1+PRs merged) + 1.0·ln(1+PRs opened)
score    = activity × 0.5^(days since last activity / 180)
risk     = permission weight / (1 + score)          admin 3 · maintain 2 · write 1.5 · triage/read 0.5
```

Flagged below **score 5.0**; suggestions ranked by **risk**, highest first.

- **Logarithmic**, because the gap between 0 and 5 commits matters and the gap between 200 and 205 does not.
- **Reviews weighted 2.5, nearly as high as commits** — this is the most consequential number here. A naive commit count flags exactly the wrong person: the senior engineer who reviews constantly, writes little code, and holds the most load-bearing access on the team. A dedicated test asserts that person is never flagged.
- **Half-life decay, not a cutoff**, because a hard 180-day boundary makes someone at 179 days safe and at 181 days flagged — indefensible in the conversation that follows.
- **Risk divides power by usage**, because a reviewer triaging a list needs "inactive *and* dangerous", not just "inactive". A dormant admin outranks a dormant reader.

Assumptions and edge cases

Handled explicitly, each with a test: empty repos (409), archived repos (**scanned, not skipped** — that access is still live, and archiving only changes the remediation advice), forks (skipped, configurable), bots, org owners, brand-new repos, single-contributor repos, custom repository roles (weighted as *write*, never as harmless).

Verified on a live organization, not only on fixtures: 7 repositories, one org-wide call returning 32 Dependabot alerts (15 critical or high), and 2 access suggestions out of 21 grants assessed — while correctly leaving alone an active admin, two org owners, a member whose read access comes from an org-wide setting, and a contributor holding no access at all. Running against real data also exposed four bugs that fixtures could not — a refused access listing rendering as "nobody has access", organization base permission mislabelled as a team grant, contributors with no access being suggested for removal, and advisory counts being silently zeroed by a second writer.

Unattributed commits are counted and reported, never guessed at. When a commit's author email is not linked to a GitHub account, matching on raw name or email would silently credit the wrong person. The live run found 6 such commits and said so.

Four access paths are stored and reported separately: direct, outside collaborator, team-inherited, and the organization's own base permission. Each has a different remediation — revoking a team grant affects every repo that team can reach; base permission is one org-wide setting — so merging them into a single "permission" column would destroy the distinction that makes the report actionable. The base-permission path was added after a live scan showed it being mislabelled as team-inherited.

Deliberately skipped: commits outside the default branch, PR reviews beyond the first 50 on a single PR, and deploy keys as an access vector. Each is noted where it applies.

Reliability

Idempotent by construction: every table keyed on natural GitHub IDs with `INSERT ... ON CONFLICT DO UPDATE` , progress committed per repo so `--resume` continues after a crash. **Verified live** — two full runs left every fact table unchanged.

Rate limits: pre-emptive pausing before the quota runs out, `Retry-After` honoured, and stored ETags so re-runs cost little. **Verified live** — the second run served 8 of 13 requests from 304s, which do not count against quota. A 100-repo org costs roughly 410 REST requests against a 5,000/hour budget, because advisories are fetched org-wide in one call rather than once per repo, and contributions use one GraphQL query per repo rather than one per member.

Production-readiness

The three things I would add next: an accept/reject feedback loop so the threshold is defensible with evidence rather than argument; tracking access *changes* between runs, since "granted admin last week, never used it" beats any snapshot; and async repo scanning, which is the obvious bottleneck past a few hundred repos.

What to look at

Run it in 30 seconds	<code>python main.py --demo</code> — no token, no org, pinned clock, reproducible output
Dashboard	docs/demo_dashboard.html — open straight from a clone
Reasoning and trade-offs	DESIGN_NOTES.md — includes AI-assistance disclosure (§11)
Execution walkthrough	RUN_REPORT.md — captured output, demo <i>and</i> live org
Tests	<code>python -m pytest -q</code> — 156 tests, ~0.3s, also run in CI on 3.11 and 3.12

What is mocked, and why: only the API responses in `--demo` mode, because private Dependabot advisories need an org with paid security features and genuinely vulnerable dependencies. The fixtures are shaped like real API payloads and run through the *same* parsing, scoring and rendering code as a live scan. Section 8 of the run report is a real scan against a live organization.